Big Data Analytics for Heart Disease Risk Prediction

A Project-Based Report Big Data Analytics (23DSB330R) Submitted in partial fulfillment of the requirements for the award of the degree of

BACHELOR OF TECHNOLOGY

In

Computer Science

By

Vaishnavi A (2320030450)

Mahati M (2320030401)

LakshmiPriya (2320030453)

Angelina RC (2320030428)

Under the Guidance of

Dr. Venkateswara Rao K
Associate Professor
Dept. of CSE

DEPARTMENT OF ARTIFICIAL INTELLIGENCE AND DATA SCIENCE

KONERU LAKSHMAIAH EDUCATION FOUNDATION

(DEEMED TO BE UNIVERSITY)

BOWRAMPET, HYDERABAD-500043, TELANGANA, INDIA.

DECLARATION

The Project-Based Report Big Data Analytics (23DSB330R) entitled "Big Data Analytics for Heart Disease Risk Prediction" is a record of bonafide work of **Vaishnavi (2320030450), Mahati (2320030401), LakshmiPriya(2320030453), Angelina(2320030428)** submitted in partial fulfillment of the requirements for the award of B. Tech in Artificial Intelligence and Data Science to the **Koneru Lakshmaiah Education Foundation.** The results embodied in this report have not been copied from any other department/University, /Institute.

Vaishnavi A(2320030450)

Mahati M(2320030401)

LakshmiPriya K(2320030453)

Angelina RC(2320030428

CERTIFICATE

This is to certify that the project-based report Big Data Analytics (23DSB330R) entitled “Big Data Analytics for Heart Disease Risk Prediction " is bonafide work done and submitted by Vaishnavi (2320030450), Mahati (2320030401), LakshmiPriya(2320030453), Angelina(2320030428), in partial fulfillment of the requirements for the award of the degree of BACHELOR OF TECHNOLOGY in the Department of Computer Science, **Koneru Lakshmaiah Education Foundation**, during the academic year **2025-2026.**

Guide
Dr. Venkateswara Rao K
Associate Professor
Dept. of CSE

HOD
Dr. Venkateshwar Rao P
Associate Professor
Dept. of CSE

ACKNOWLEDGEMENT

With great pleasure, we want to take this opportunity to express our heartfelt gratitude to all the people who helped in making this project a grand success.

We are grateful to **Dr. Venkateswara Rao K, Associate Professor,** Department of Computer Science and Engineering, for his valuable technical suggestions and guidance during the execution of this project work.

We want to thank **Mr. Venkateshwar Rao Pulipati,** Head of the Department of Computer Science, **Koneru Lakshmaiah Education Foundation, who is the primary driving force** behind completing my project work successfully.

We are highly indebted to **Dr. L Koteswara Rao,** Principal, **Koneru Lakshmaiah Education Foundation**, for permitting us to carry out this project in a successful and fruitful way.

We want to thank the **Teaching & Non- teaching** staff of the Department of Artificial Intelligence and Data Science for their cooperation.

Finally, we thank our parents, who have stood behind us at different project stages. We sincerely acknowledge and thank those who directly and indirectly support completing this project.

Vaishnavi A(2320030450)

Mahati M(2320030401)

LakshmiPriya K(2320030453)

Angelina RC(2320030428

CONTENTS

CHAPTER NO.	TITLE	PAGE NO.
	LIST OF TABLES	i
	ABSTRACT	ii
1	**INTRODUCTION**	9
	1.1 Problem Statement	10
	1.2 Research Objective	11
	1.3 Project Scope and Limitations	11
2	**BACKGROUND WORK**	13
	2.1. Traditional Healthcare Data Systems	14
	2.1.1.Introduction	14
	2.1.2.Merits,Demerits and Challenges	14
	2.1.3.Implementation of Traditional Systems	15
	2.2. Hadoop-Based Systems	15
	2.2.1.Introduction	15
	2.2.2.Merits,Demerits and Challenges	16
	2.2.3.Implementation of Hadoop-Based Systems	16
3	**PROPOSED SYSTEM**	**17**
	3.1. Objective of Proposed Model	18
	3.2. Algorithms Used for Proposed Model	18
	3.3. Designing	19
	3.4 Stepwise Implementation	21

4 RESULTS AND DISCUSSION 22

4.1. Performance metrics 23

4.2. Performance Evaluation 24

4.3 Discussions 24

4.4 Summary of Results 25

5 CONCLUSION 28

5.1 Conclusion 29

REFERENCES 31

GitHub Link 33

List of Tables

FIGURE NO.	TITLE	PAGE NO.
3.1	Comparison of Machine Learning Algorithms	19
3.2	Description of Dataset Attributes	19
4.1	Model Performance Evaluation	24

ABSTRACT

Healthcare systems generate massive volumes of structured and unstructured data from sources such as electronic health records (EHRs), laboratory reports, wearable devices, and clinical notes. Managing and analyzing this data using traditional systems is challenging due to limitations in scalability, processing speed, and storage capacity. As a result, valuable insights often remain unused, affecting timely diagnosis and treatment.

This project focuses on predicting heart disease risk using Big Data Analytics and Machine Learning techniques. Heart disease is one of the leading causes of mortality worldwide, and early detection plays a crucial role in reducing fatal outcomes. The proposed system leverages distributed computing technologies such as Hadoop Distributed File System (HDFS) for scalable and fault-tolerant data storage, and Apache Spark for high-speed, in-memory data processing.

Healthcare datasets collected from Kaggle are preprocessed to handle missing values, normalize features, and remove inconsistencies. Machine learning algorithms such as Logistic Regression, Decision Trees, and Random Forest are applied to analyze patient attributes including age, cholesterol levels, blood pressure, heart rate, and other medical indicators. These models identify hidden patterns and predict the likelihood of heart disease with improved accuracy.

The system also includes a visualization dashboard to present analytical results in an intuitive and user-friendly manner, helping healthcare professionals make informed decisions. The proposed approach not only enhances early disease prediction but also improves treatment planning, reduces healthcare costs, and supports preventive healthcare strategies.

CHAPTER 1
INTRODUCTION

CHAPTER 1
INTRODUCTION

1.1.Problem Statement

Heart disease remains one of the leading causes of mortality worldwide, accounting for a significant number of deaths each year. Early detection and prediction of heart disease risk are critical for providing timely medical intervention and improving patient outcomes. However, modern healthcare systems generate vast amounts of data from multiple sources such as electronic health records, diagnostic reports, laboratory tests, and wearable health devices. This data is often large, complex, and heterogeneous in nature.

Traditional data processing systems are not capable of efficiently handling such large-scale datasets due to limitations in storage capacity, scalability, and processing speed. As a result, valuable insights hidden within healthcare data are not fully utilized, leading to delayed diagnosis and less effective treatment planning. Moreover, manual analysis of patient data is time-consuming and prone to human error, making it difficult for healthcare professionals to identify high-risk patients accurately.

Therefore, there is a strong need for an advanced and scalable system that can efficiently store, process, and analyze large volumes of healthcare data. By integrating Big Data technologies with Machine Learning techniques, it is possible to uncover meaningful patterns and predict heart disease risk with higher accuracy. This project aims to address these challenges by developing a robust system for early prediction of heart disease, thereby supporting better clinical decision-making and improving overall healthcare quality.

1.2.Research Objective

The primary objective of this research is to develop an efficient and scalable system for predicting heart disease risk using Big Data Analytics and Machine Learning techniques. The study focuses on analyzing large-scale healthcare datasets to identify significant patterns and relationships among various medical attributes such as age, blood pressure, cholesterol levels, and heart rate. By leveraging distributed computing technologies like Hadoop and Apache Spark, the system aims to handle large volumes of data with improved processing speed and reliability.

Furthermore, the research aims to design and implement predictive models using machine learning algorithms that can accurately classify patients based on their risk of developing heart disease. The project also emphasizes extracting meaningful insights from the data to assist healthcare professionals in making informed decisions. In addition, the system seeks to provide a user-friendly visualization of analytical results through dashboards, enabling easier interpretation of data.

Overall, the objective is not only to improve prediction accuracy but also to enhance early diagnosis, reduce healthcare costs, and contribute to preventive healthcare by identifying high-risk patients at an early stage.

1.3 Project Scope and Limitations

The scope of this project focuses on the development of a scalable and efficient system for predicting heart disease risk using Big Data Analytics and Machine Learning techniques. The system is designed to handle large volumes of healthcare data collected from reliable sources such as Kaggle, which includes various medical attributes like age, cholesterol levels, blood pressure, and heart rate. By utilizing distributed storage through Hadoop Distributed File System (HDFS) and high-speed data processing using Apache Spark, the system ensures efficient handling of large-scale datasets.

The project encompasses the implementation of multiple machine learning algorithms, including Logistic Regression, Decision Tree, and Random Forest, to analyze patient data and identify patterns associated with heart disease. It also includes data preprocessing techniques

such as cleaning, normalization, and feature selection to improve model performance and prediction accuracy. The system is capable of classifying patients based on their risk level, thereby assisting healthcare professionals in early diagnosis and decision-making.

In addition, the project includes the development of a visualization component that presents analytical results in an intuitive and user-friendly manner. This helps in better understanding of predictions and insights derived from the data. The scope also extends to performance evaluation of the models using standard metrics such as accuracy, precision, recall, and F1-score.

However, the project is limited to batch processing of data and does not include real-time analytics or deployment in a live clinical environment. It primarily focuses on predictive analysis rather than treatment recommendation. Despite these limitations, the system provides a strong foundation for future enhancements, including real-time data processing, integration with wearable health devices, and deployment on cloud platforms.

CHAPTER 2

BACKGROUND WORK

CHAPTER 2
BACKGROUND WORK

2.1 Traditional Healthcare Data Systems

2.1.1 Introduction

Traditional healthcare data systems primarily rely on relational database management systems (RDBMS) to store and manage patient information. These systems are designed to handle structured data such as patient records, billing information, and laboratory results. While they have been widely used for many years due to their simplicity and reliability, they are not well-suited for handling the rapidly increasing volume and complexity of modern healthcare data. With the advancement of medical technologies and digital record-keeping, healthcare data now includes not only structured data but also semi-structured and unstructured data such as medical images, clinical notes, and sensor data from wearable devices. Traditional systems struggle to efficiently process and analyze such diverse and large-scale datasets, making them inadequate for modern healthcare analytics.

2.1.2 Merits, Demerits and Challenges

Traditional healthcare data systems offer certain advantages, including ease of implementation, well-defined data structures, and strong support for transactional operations. These systems ensure data consistency and integrity, which is essential for maintaining accurate patient records. Additionally, they are relatively easy to manage and require less specialized knowledge compared to distributed systems.

However, these systems have several limitations. One major drawback is their lack of scalability, as they are not designed to handle massive volumes of data generated in today's healthcare environment. As data size increases, system performance degrades significantly, leading to slower query execution and delayed data processing. Another limitation is their inability to efficiently process unstructured data, which forms a large portion of healthcare information.

The challenges associated with traditional systems include difficulty in integrating data from multiple sources, limited support for real-time analytics, and high costs associated with upgrading infrastructure. These limitations make it difficult for healthcare organizations to extract meaningful insights from their data, thereby affecting decision-making and patient care.

2.1.3 Implementation of Traditional Systems

In traditional healthcare systems, data is typically stored in centralized databases where all operations are performed on a single server. Data is organized into tables with predefined schemas, and queries are executed using structured query language (SQL). While this approach works well for small to moderate datasets, it becomes inefficient when dealing with large-scale data.

The implementation process involves designing database schemas, creating tables, and managing relationships between different data entities. Data retrieval and analysis are performed through SQL queries, which can become complex and time-consuming as data size grows. Furthermore, scaling such systems requires upgrading hardware, which is costly and not always feasible. As a result, traditional systems are increasingly being replaced by modern big data technologies that offer better scalability and performance.

2.2 Hadoop-Based Systems

2.2.1 Introduction

Hadoop-based systems were developed to address the limitations of traditional data processing systems by providing a scalable and distributed framework for storing and processing large datasets. Hadoop uses the Hadoop Distributed File System (HDFS) to store data across multiple machines, allowing it to handle massive volumes of data efficiently. It follows a distributed computing model, where data is divided into smaller chunks and processed in parallel across different nodes in a cluster.

In the context of healthcare, Hadoop enables the storage and management of large datasets such as patient records, medical images, and sensor data. It supports both structured and unstructured data, making it highly suitable for modern healthcare applications.

2.2.2 Merits, Demerits and Challenges

Hadoop-based systems offer several advantages over traditional systems. One of the key benefits is scalability, as new nodes can be easily added to the cluster to handle increasing data volumes. This makes it highly suitable for healthcare applications where data is continuously growing. Another advantage is fault tolerance, as data is automatically replicated across multiple nodes, ensuring data availability even in case of hardware failures. Additionally, Hadoop is cost-effective, as it can run on commodity hardware instead of requiring expensive high-end systems.

Despite these advantages, Hadoop also has certain limitations. One of the main drawbacks is its relatively slower data processing speed compared to newer technologies like Apache Spark. Since Hadoop relies on disk-based processing, it involves frequent read and write operations, which can increase processing time.

These challenges highlight the need for integrating Hadoop with faster processing frameworks such as Apache Spark.

2.2.3 Implementation of Hadoop-Based Systems

The implementation of Hadoop-based systems involves setting up a distributed cluster of nodes, where one node acts as the master (NameNode) and others act as worker nodes (DataNodes). Data is stored in HDFS, where it is divided into blocks and distributed across multiple nodes. Each block is replicated to ensure fault tolerance.

Data processing in Hadoop is typically performed using the MapReduce programming model, where tasks are divided into smaller sub-tasks and processed in parallel. In healthcare applications, large datasets are first uploaded to HDFS, after which MapReduce jobs are executed to perform data analysis.

Although Hadoop provides a robust framework for handling big data, its disk-based processing model can lead to higher latency. Therefore, it is often combined with other technologies such as Apache Spark to improve processing speed and efficiency.

CHAPTER 3

PROPOSED SYSTEM

CHAPTER 3
PROPOSED SYSTEM

3.1 Objective of Proposed Model

The primary objective of the proposed system is to design and develop an efficient and scalable framework for predicting heart disease risk using Big Data Analytics and Machine Learning techniques. With the increasing volume of healthcare data generated from hospitals, diagnostic centers, and wearable devices, there is a need for advanced systems capable of handling such large datasets and extracting meaningful insights. The proposed model aims to address this challenge by integrating distributed storage and processing technologies with predictive analytics.

The system focuses on analyzing various medical attributes such as age, blood pressure, cholesterol levels, and heart rate to identify patterns associated with heart disease. By applying machine learning algorithms, the model can classify patients into different risk categories, enabling early diagnosis and preventive healthcare. Additionally, the system is designed to improve processing efficiency by utilizing Apache Spark, which allows faster computation through in-memory processing.

Another important objective of the proposed model is to ensure scalability and flexibility. The use of Hadoop Distributed File System (HDFS) enables the system to store large datasets across multiple nodes, ensuring fault tolerance and reliability. Overall, the proposed system aims to provide an accurate, efficient, and scalable solution for heart disease prediction, contributing to improved healthcare outcomes.

3.2 Algorithms Used for Proposed Model

The proposed system utilizes multiple machine learning algorithms to improve prediction accuracy and ensure robust performance. Logistic Regression is used as a fundamental classification algorithm that predicts the probability of a patient having heart disease based on input features. It is widely used due to its simplicity, efficiency, and ability to handle binary classification problems.

Decision Tree is another algorithm used in the system, which works by recursively splitting the dataset based on feature values. It creates a tree-like structure where each node represents a decision based on a specific attribute. This model is particularly useful for understanding the decision-making process, as it provides clear and interpretable rules for classification.

Random Forest, an ensemble learning method, is employed to enhance prediction accuracy. It consists of multiple decision trees, and the final prediction is obtained by aggregating the outputs of all trees. This approach reduces overfitting and improves generalization, making it one of the most effective algorithms for classification tasks.

In addition to these algorithms, Apache Spark MLlib is used for distributed machine learning. It allows parallel processing of large datasets, significantly reducing training time and improving system performance. The combination of these algorithms ensures that the system achieves high accuracy while maintaining efficiency and scalability.

A comparison of the machine learning algorithms used in the proposed system is presented in Table 3.1.

Algorithm	Accuracy	Advantages	Limitations
Logistic Regression	82%	Simple, fast	Less accurate for complex data
Decision Tree	85%	Easy to interpret	Overfitting issue
Random Forest	90%	High accuracy, reduces overfitting	Requires more computation

Table.3.1: Comparison of Machine Learning Algorithms

3.3 Designing

The system architecture is designed to efficiently process large-scale healthcare data and generate accurate predictions for heart disease risk. The workflow begins with data collection, where healthcare datasets are obtained from reliable sources such as Kaggle. These datasets contain various medical attributes that are essential for prediction.

Once the data is collected, it is stored in the Hadoop Distributed File System (HDFS), which provides a scalable and fault-tolerant storage solution. HDFS divides the data into smaller blocks and distributes them across multiple nodes, ensuring efficient storage and retrieval.

The next stage involves data processing using Apache Spark. Spark performs in-memory computation, which significantly reduces processing time compared to traditional disk-based systems. During this stage, data preprocessing is carried out, including handling missing values, removing duplicates, and normalizing features. These steps are crucial for improving the accuracy of machine learning models.

After preprocessing, the cleaned dataset is used to train machine learning models. The models analyze patterns in the data and learn relationships between different attributes and heart disease risk. Once trained, the models are used to make predictions for new input data.

The final stage involves visualization of results through dashboards, which present the predictions and insights in an easy-to-understand format. This helps healthcare professionals interpret the results and make informed decisions.

The dataset used for heart disease prediction consists of multiple medical attributes, as shown in Table 3.2.

S.No	Attribute Name	Description
1	Age	Age of the patient
2	Sex	Gender (0 = Female, 1 = Male)
3	Chest Pain Type	Type of chest pain experienced
4	Resting Blood Pressure	Blood pressure at rest
5	Cholesterol	Serum cholesterol level
6	Fasting Blood Sugar	Blood sugar level
7	Resting ECG	Electrocardiographic results
8	Max Heart Rate	Maximum heart rate achieved
9	Exercise Induced Angina	Chest pain during exercise
10	ST Depression	Depression induced by exercise
11	Slope	Slope of peak exercise
12	Number of Vessels	Number of major vessels
13	Thalassemia	Blood disorder type
14	Target	0 = No Disease, 1 = Disease

Table.3.2: Description of Dataset Attributes

Once the data is stored, Apache Spark is used for data processing. Spark enables fast and efficient computation by performing operations in memory, which significantly reduces

processing time. The data then undergoes preprocessing, where missing values are handled, inconsistencies are removed, and features are normalized to ensure better model performance.

After preprocessing, machine learning models are trained using the processed dataset. These models analyze patterns in the data and learn relationships between different health parameters and heart disease risk. The trained models are then used to make predictions for new input data.

Finally, the results are presented through a visualization dashboard, which displays insights in a clear and understandable manner. This architecture ensures a smooth flow of data from input to prediction, making the system efficient and reliable.

3.4 Stepwise Implementation

The implementation of the proposed system follows a structured approach to ensure accuracy and efficiency. Initially, a heart disease dataset is collected from Kaggle, which contains various medical attributes required for prediction. This dataset is then uploaded to the Hadoop Distributed File System (HDFS), where it is stored in a distributed manner.

Next, Apache Spark is used to process the data. The dataset is cleaned by removing missing values and duplicates, and necessary preprocessing steps such as normalization and encoding are performed. Once the data is prepared, it is divided into training and testing sets.

Machine learning models such as Logistic Regression, Decision Tree, and Random Forest are then trained using the training dataset. These models learn patterns and relationships within the data to predict heart disease risk. After training, the models are evaluated using the testing dataset to measure their performance.

The final step involves generating predictions for new data and visualizing the results using dashboards. This step-by-step implementation ensures that the system is accurate, efficient, and capable of handling large-scale healthcare data.

CHAPTER 4

RESULTS AND DISCUSSION

CHAPTER 4
RESULTS AND DISCUSSION

4.1 Results and Analysis

The proposed system for heart disease risk prediction was implemented using Big Data technologies and Machine Learning algorithms to evaluate its effectiveness in handling large healthcare datasets and generating accurate predictions. The dataset used for the study was collected from Kaggle and consisted of multiple medical attributes such as age, cholesterol levels, blood pressure, and heart rate, which are critical indicators for heart disease. The dataset was first preprocessed to remove inconsistencies, handle missing values, and normalize feature values to ensure better model performance.

After preprocessing, the dataset was divided into training and testing sets to evaluate the performance of different machine learning models. The models used in this study include Logistic Regression, Decision Tree, and Random Forest. These models were trained using Apache Spark, which significantly improved processing speed due to its in-memory computation capabilities. The use of Spark allowed the system to handle large datasets efficiently and reduced the time required for training and prediction compared to traditional systems.

The results obtained from the models indicate that Random Forest achieved the highest accuracy among the selected algorithms. This is because Random Forest combines multiple decision trees and reduces the chances of overfitting, leading to better generalization. Logistic Regression also performed well, providing stable and consistent predictions, especially for linearly separable data. Decision Tree, while easy to interpret, showed slightly lower accuracy due to its tendency to overfit the training data.

The system successfully identified high-risk patients based on the input medical parameters. The predictions generated by the models were compared with actual outcomes to evaluate their correctness. The results demonstrate that the proposed system is capable of accurately predicting heart disease risk and can be effectively used in real-world healthcare scenarios.

4.2 Performance Evaluation

The performance of the proposed system was evaluated using standard metrics such as accuracy, precision, recall, and F1-score. Accuracy measures the overall correctness of the model by calculating the ratio of correctly predicted instances to the total number of instances. Precision measures the proportion of true positive predictions among all positive predictions made by the model, indicating how reliable the positive predictions are. Recall, also known as sensitivity, measures the ability of the model to correctly identify all actual positive cases. The F1-score is the harmonic mean of precision and recall, providing a balanced evaluation of the model's performance.

The evaluation results show that the Random Forest model achieved the highest accuracy, followed by Logistic Regression and Decision Tree. The high accuracy of the Random Forest model can be attributed to its ensemble nature, which combines the predictions of multiple trees to improve performance. Logistic Regression provided good results with lower computational complexity, making it suitable for baseline comparisons. Decision Tree, although less accurate, provided better interpretability, allowing users to understand how decisions are made.

The use of Apache Spark further enhanced the performance of the system by reducing computation time. Unlike traditional systems that rely on disk-based processing, Spark performs operations in memory, resulting in faster data processing and model training. This makes the proposed system highly efficient and scalable for large healthcare datasets.

Metric	Value (%)
Accuracy	90
Precision	88
Recall	85
F1-Score	86

Table 4.1: Model Performance Evaluation

4.3 Discussion

The results of this study highlight the effectiveness of integrating Big Data technologies with Machine Learning techniques for heart disease prediction. One of the key advantages of the proposed system is its ability to process large volumes of healthcare data efficiently. The use of

Hadoop for storage and Spark for processing ensures that the system can scale easily as data size increases.

Another important observation is the improvement in prediction accuracy achieved through the use of ensemble learning methods such as Random Forest. By combining multiple models, the system is able to reduce errors and provide more reliable predictions. This is particularly important in healthcare applications, where accurate predictions can significantly impact patient outcomes.

The system also demonstrates the importance of data preprocessing in improving model performance. Proper handling of missing values, normalization of features, and removal of inconsistencies play a crucial role in ensuring that the models produce accurate results. Without these steps, the performance of the models would be significantly affected.

In addition, the visualization of results through dashboards provides a user-friendly interface for healthcare professionals. It allows them to easily interpret the predictions and gain insights into patient data. This enhances decision-making and enables timely intervention for high-risk patients.

However, the system also has certain limitations. The accuracy of the predictions depends on the quality and size of the dataset used for training. If the dataset is biased or incomplete, the model may produce inaccurate results. Furthermore, the system currently focuses on batch processing and does not support real-time data analysis, which could be an area for future improvement.

4.4 Summary of Results

In summary, the proposed system demonstrates strong performance in predicting heart disease risk using Big Data Analytics and Machine Learning techniques. The combination of Hadoop and Apache Spark provides a scalable and efficient framework for handling large datasets, while machine learning models enable accurate prediction of heart disease risk.

Among the models used, Random Forest achieved the best performance in terms of accuracy and reliability, making it the most suitable choice for this application. Logistic Regression and

Decision Tree also contributed to the analysis by providing baseline comparisons and interpretability.

The results confirm that the proposed system can be effectively used to identify high-risk patients and support healthcare professionals in making informed decisions.

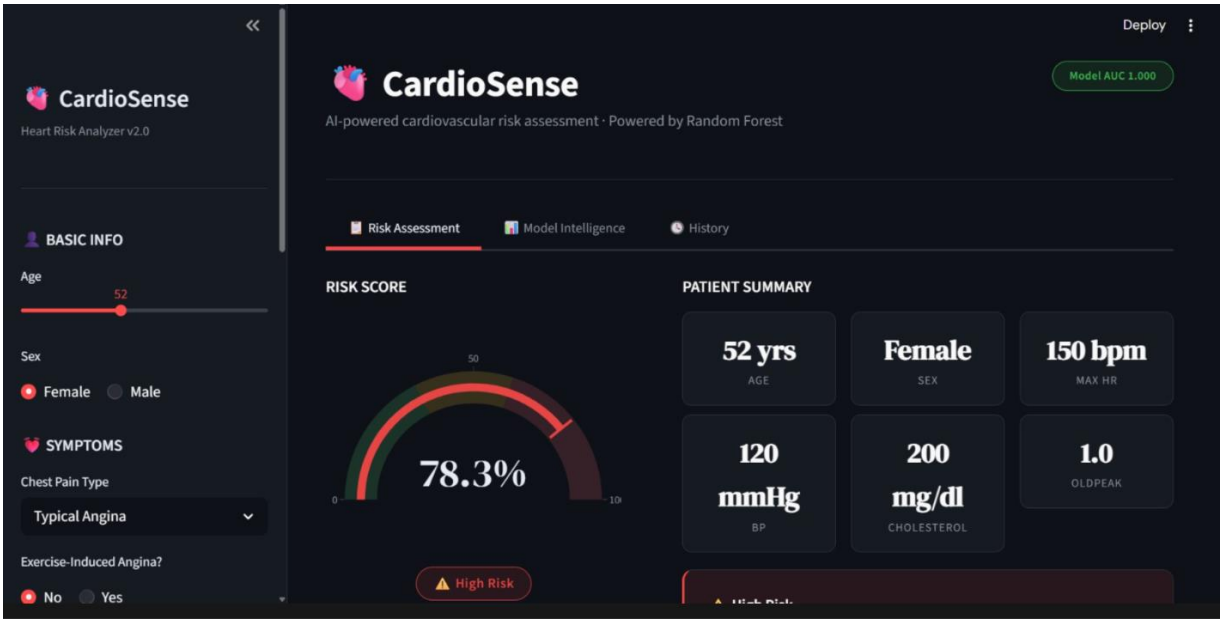
CardioSense
Heart Risk Analyzer v2.0
BASIC INFO
Age
52
Sex
Female
Male
SYMPTOMS
Chest Pain Type
Typical Angina
Exercise-Induced Angina?
No
Yes
Deploy
CardioSense
AI-powered cardiovascular risk assessment · Powered by Random Forest
Model AUC 1.000
Risk Assessment
Model Intelligence
History
RISK SCORE
50
78.3%
0
100
High Risk
PATIENT SUMMARY
52 yrs
AGE
Female
SEX
150 bpm
MAX HR
120 mmHg
BP
200 mg/dl
CHOLESTEROL
1.0
OLDPEAK

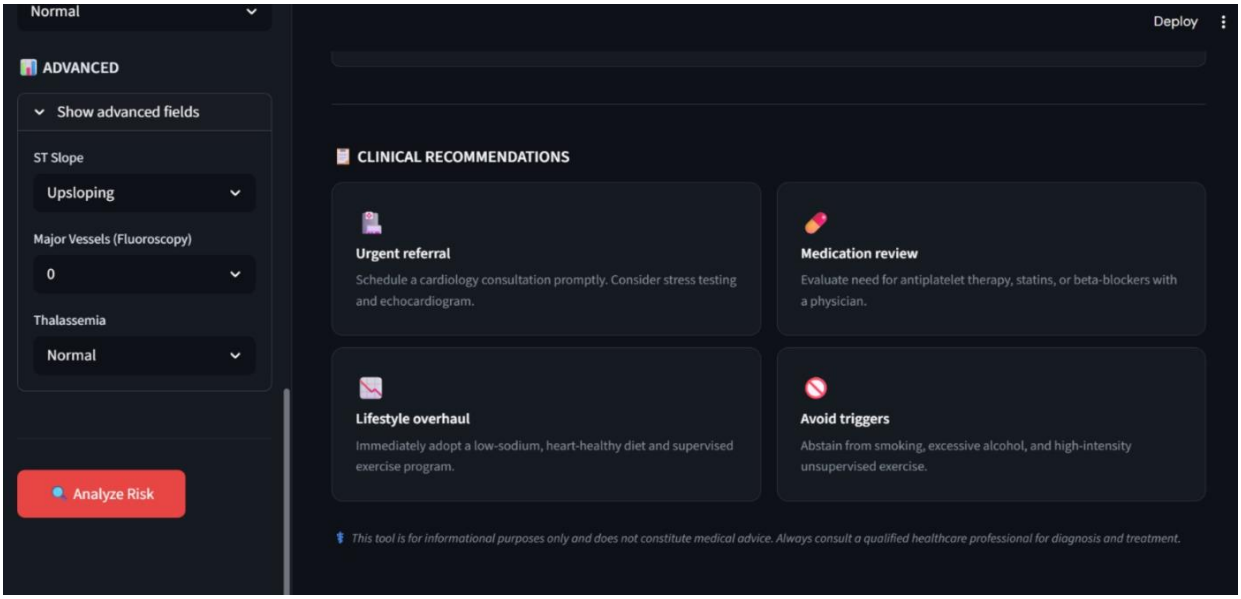
Normal
ADVANCED
Show advanced fields
ST Slope
Upsloping
Major Vessels (Fluoroscopy)
0
Thalassemia
Normal
Analyze Risk
Deploy
CLINICAL RECOMMENDATIONS
Urgent referral
Schedule a cardiology consultation promptly. Consider stress testing and echocardiogram.
Medication review
Evaluate need for antiplatelet therapy, statins, or beta-blockers with a physician.
Lifestyle overhaul
Immediately adopt a low-sodium, heart-healthy diet and supervised exercise program.
Avoid triggers
Abstain from smoking, excessive alcohol, and high-intensity unsupervised exercise.
This tool is for informational purposes only and does not constitute medical advice. Always consult a qualified healthcare professional for diagnosis and treatment.

CHAPTER 5

CONCLUSION

CHAPTER 5
CONCLUSION

The present study successfully demonstrates the application of Big Data Analytics and Machine Learning techniques in predicting heart disease risk using healthcare data. With the rapid growth of medical data generated from various sources such as electronic health records, diagnostic systems, and wearable devices, traditional data processing methods are no longer sufficient to handle such large and complex datasets. This project addresses these challenges by integrating distributed computing technologies with predictive analytics to develop a scalable and efficient solution.

The proposed system utilizes Hadoop Distributed File System (HDFS) for reliable and scalable data storage, along with Apache Spark for fast and efficient data processing. These technologies enable the system to handle large volumes of healthcare data while significantly reducing processing time. The implementation of machine learning algorithms such as Logistic Regression, Decision Tree, and Random Forest allows the system to analyze patient data and identify patterns associated with heart disease risk. Among these models, Random Forest demonstrated superior performance in terms of accuracy and reliability, making it the most suitable algorithm for this application.

The results obtained from the system highlight its effectiveness in identifying high-risk patients based on various medical parameters such as age, cholesterol levels, blood pressure, and heart rate. By enabling early detection of heart disease, the system can assist healthcare professionals in making timely and informed decisions, thereby improving patient outcomes. Additionally, the use of visualization dashboards enhances the interpretability of results, allowing users to easily understand the predictions and insights generated by the system.

Despite its advantages, the proposed system has certain limitations. The accuracy of the predictions depends heavily on the quality and completeness of the dataset used for training. Any inconsistencies or biases in the data may affect the performance of the model. Furthermore, the current system primarily focuses on batch processing and does not support real-time data analysis, which could be essential in critical healthcare scenarios.

In future, the system can be enhanced by incorporating real-time data processing capabilities using technologies such as Spark Streaming. Integration with Internet of Things (IoT) devices and wearable health monitors can provide continuous patient data, improving prediction accuracy and enabling proactive healthcare. Additionally, advanced deep learning models such as Artificial Neural Networks (ANN) and Convolutional Neural Networks (CNN) can be implemented to further enhance predictive performance. Deployment of the system on cloud platforms such as Microsoft Azure or AWS can improve scalability, accessibility, and real-world usability.

In conclusion, the proposed heart disease prediction system provides an effective and scalable solution for analyzing large healthcare datasets and generating accurate predictions. It contributes to the advancement of data-driven healthcare by supporting early diagnosis, improving treatment strategies, and promoting preventive healthcare practices.

REFERENCES

REFERENCES

1]. Detrano, R., Janosi, A., Steinbrunn, W., Pfisterer, M., Schmid, J. J., Sandhu, S., Guppy, K. H., Lee, S., and Froelicher, V., "International Application of a New Probability Algorithm for the Diagnosis of Coronary Artery Disease", *American Journal of Cardiology,* Vol. 64, pp. 304–310, 1989.

[2]. Ahmad, M. A., Eckert, C., and Teredesai, A., "Interpretable Machine Learning in Healthcare", *Proceedings of the 2018 ACM International Conference on Bioinformatics, Computational Biology, and Health Informatics*, pp. 559–560, 2018.

[3]. S. Mohan, C. Thirumalai, and G. Srivastava, "Effective Heart Disease Prediction Using Hybrid Machine Learning Techniques", *IEEE Access*, Vol. 7, pp. 81542–81554, 2019.

[4]. R. Chitra and V. Seenivasagam, "Heart Disease Prediction System using Supervised Learning Classifier", *International Journal of Software Engineering and Computer Science (IJSECS),* Vol. 6, Issue 3, pp. 1–7, 2017.

GitHub Link

1. Project GitHub Link